

TITANIC CLASSIFICATION

Build a predictive model to determine the likelihood of survival for passengers on the Titanic using data science techniques in Python.

```
In [1]: # Importing Libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: # Loading the Dataset
titanic = pd.read_csv('Titanic-Dataset.csv')
titanic
```

```
Out[3]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.4500
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.7500

891 rows × 12 columns



```
In [4]: # Reading first 5 rows
titanic.head()
```

```
Out[4]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	N
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	N
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C1
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	N

```
In [5]: # Reading Last 5 rows
titanic.tail()
```

```
Out[5]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.00	NaN
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.00	B42
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.45	NaN
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.00	C148
890	891	0	3	Dooley, Mr. Patrick	male	32.0	0	0	370376	7.75	NaN

```
In [6]: # Showing no. of rows and columns of dataset
titanic.shape
```

```
Out[6]: (891, 12)
```

```
In [7]: # checking for columns
titanic.columns
```

```
Out[7]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
              'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
              dtype='object')
```

```
In [8]: # Checking for data types
titanic.dtypes
```

```
Out[8]: PassengerId    int64
Survived              int64
Pclass                int64
Name                  object
Sex                   object
Age                   float64
SibSp                 int64
Parch                 int64
Ticket                object
Fare                  float64
Cabin                 object
Embarked              object
dtype: object
```

```
In [9]: # checking for duplicated values
titanic.duplicated().sum()
```

```
Out[9]: 0
```

```
In [10]: # checking for null values
nv = titanic.isna().sum().sort_values(ascending=False)
nv = nv[nv>0]
nv
```

```
Out[10]: Cabin         687
Age              177
Embarked         2
dtype: int64
```

```
In [11]: # Cheeeking what percentage column contain missing values
titanic.isnull().sum().sort_values(ascending=False)*100/len(titanic)
```

```
Out[11]: Cabin          77.104377
Age             19.865320
Embarked        0.224467
PassengerId     0.000000
Survived        0.000000
Pclass          0.000000
Name            0.000000
Sex             0.000000
SibSp           0.000000
Parch           0.000000
Ticket          0.000000
Fare            0.000000
dtype: float64
```

```
In [12]: # Since Cabin Column has more than 75 % null values .So , we will drop this col
titanic.drop(columns = 'Cabin', axis = 1, inplace = True)
titanic.columns
```

```
Out[12]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
               'Parch', 'Ticket', 'Fare', 'Embarked'],
              dtype='object')
```

```
In [13]: # Filling Null Values in Age column with mean values of age column
titanic['Age'].fillna(titanic['Age'].mean(),inplace=True)

# filling null values in Embarked Column with mode values of embarked column
titanic['Embarked'].fillna(titanic['Embarked'].mode()[0],inplace=True)
```

```
In [14]: # checking for null values
titanic.isna().sum()
```

```
Out[14]: PassengerId    0
Survived              0
Pclass               0
Name                 0
Sex                  0
Age                  0
SibSp                0
Parch                0
Ticket               0
Fare                 0
Embarked             0
dtype: int64
```

```
In [15]: # Finding no. of unique values in each column of dataset
titanic[['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
        'Parch', 'Ticket', 'Fare', 'Embarked']].nunique().sort_values()
```

```
Out[15]: Survived      2
Sex              2
Pclass          3
Embarked        3
SibSp           7
Parch           7
Age            89
Fare          248
Ticket         681
PassengerId    891
Name           891
dtype: int64
```

```
In [17]: titanic['Survived'].unique()
```

```
Out[17]: array([0, 1], dtype=int64)
```

```
In [18]: titanic['Sex'].unique()
```

```
Out[18]: array(['male', 'female'], dtype=object)
```

```
In [19]: titanic['Pclass'].unique()
```

```
Out[19]: array([3, 1, 2], dtype=int64)
```

```
In [20]: titanic['SibSp'].unique()
```

```
Out[20]: array([1, 0, 3, 4, 2, 5, 8], dtype=int64)
```

```
In [21]: titanic['Parch'].unique()
```

```
Out[21]: array([0, 1, 2, 5, 3, 4, 6], dtype=int64)
```

```
In [22]: titanic['Embarked'].unique()
```

```
Out[22]: array(['S', 'C', 'Q'], dtype=object)
```

```
In [23]: titanic.drop(columns=['PassengerId', 'Name', 'Ticket'],axis=1,inplace=True)
titanic.columns
```

```
Out[23]: Index(['Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare',
               'Embarked'],
              dtype='object')
```

```
In [24]: # Showing information about the dataset
titanic.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Survived    891 non-null    int64
 1   Pclass      891 non-null    int64
 2   Sex         891 non-null    object
 3   Age         891 non-null    float64
 4   SibSp       891 non-null    int64
 5   Parch       891 non-null    int64
 6   Fare        891 non-null    float64
 7   Embarked    891 non-null    object
dtypes: float64(2), int64(4), object(2)
memory usage: 55.8+ KB
```

```
In [25]: # showing info. about numerical columns
titanic.describe()
```

```
Out[25]:
```

	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	13.002015	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	22.000000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	29.699118	0.000000	0.000000	14.454200
75%	1.000000	3.000000	35.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```
In [26]: # showing info. about categorical columns
titanic.describe(include='O')
```


```
Out[26]:
```

	Sex	Embarked
count	891	891
unique	2	3
top	male	S
freq	577	646


```
In [27]: ##Sex Column  
d1 = titanic['Sex'].value_counts()  
d1
```

```
Out[27]: male      577  
female    314  
Name: Sex, dtype: int64
```


```
In [28]: # Plotting Count plot for sex column  
sns.countplot(x=titanic['Sex'])  
plt.show()
```





```
In [29]: # Plotting Percentage Distribution of Sex Column  
plt.figure(figsize=(5,5))  
plt.pie(d1.values,labels=d1.index,autopct='%.2f%%')  
plt.legend()  
plt.show()
```




```
In [30]: # Showing Distribution of Sex Column Survived Wise  
sns.countplot(x=titanic['Sex'],hue=titanic['Survived']) # In Sex (0 represents  
plt.show()
```




```
In [31]: # Showing Distribution of Embarked Sex wise  
sns.countplot(x=titanic['Embarked'],hue=titanic['Sex'])  
plt.show()
```




```
In [32]: # Plotting CountPlot for Pclass Column  
sns.countplot(x=titanic['Pclass'])  
plt.show()
```



```
In [33]: # Showing Distribution of Pclass Sex wise  
sns.countplot(x=titanic['Pclass'],hue=titanic['Sex'])  
plt.show()
```



```
In [35]: # Age Distribution  
sns.kdeplot(x=titanic['Age'])  
plt.show()
```




```
In [36]: # Plotting CountPlot for Survived Column  
print(titanic['Survived'].value_counts())  
sns.countplot(x=titanic['Survived'])  
plt.show()
```


0 549

1 342


Name: Survived, dtype: int64




```
In [37]: # Showing Distribution of Parch Survived Wise  
sns.countplot(x=titanic['Parch'],hue=titanic['Survived'])  
plt.show()
```





```
In [38]: # Showing Distribution of SibSp Survived Wise  
sns.countplot(x=titanic['SibSp'],hue=titanic['Survived'])  
plt.show()
```



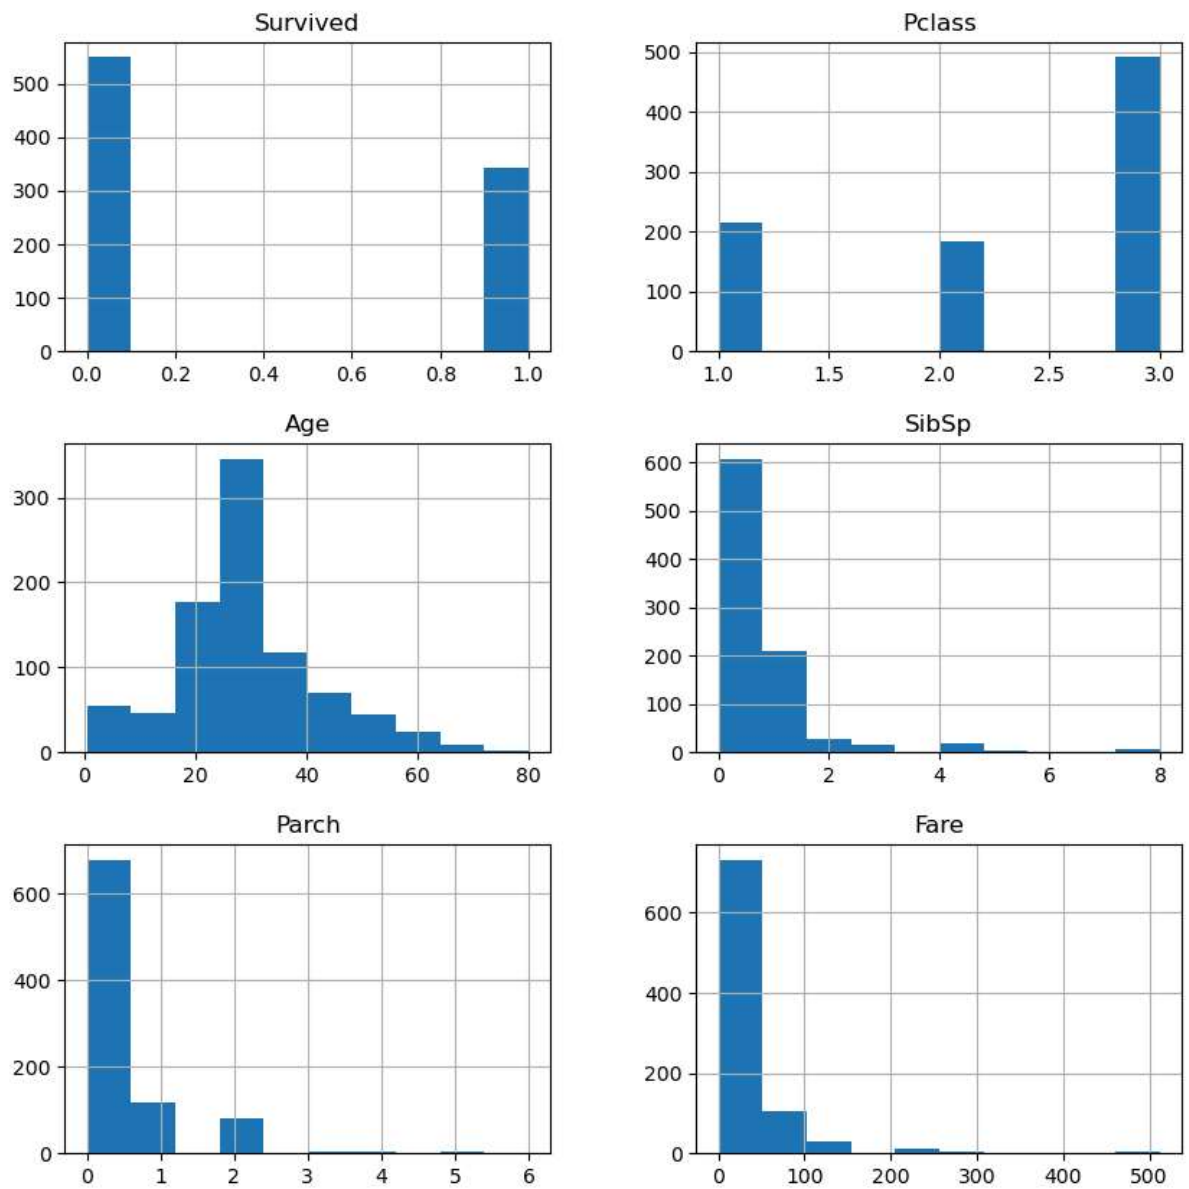
```
In [39]: # Showing Distribution of Embarked Survived wise  
sns.countplot(x=titanic['Embarked'],hue=titanic['Survived'])  
plt.show()
```




```
In [40]: # Showing Distribution of Age Survived Wise  
sns.kdeplot(x=titanic['Age'],hue=titanic['Survived'])  
plt.show()
```



```
In [41]: # Plotting Histplot for Dataset  
titanic.hist(figsize=(10,10))  
plt.show()
```



```
In [42]: # Plotting Boxplot for dataset
# Checking for outliers
sns.boxplot(titanic)
plt.show()
```




```
In [43]: # showing Correlation
titanic.corr()
```

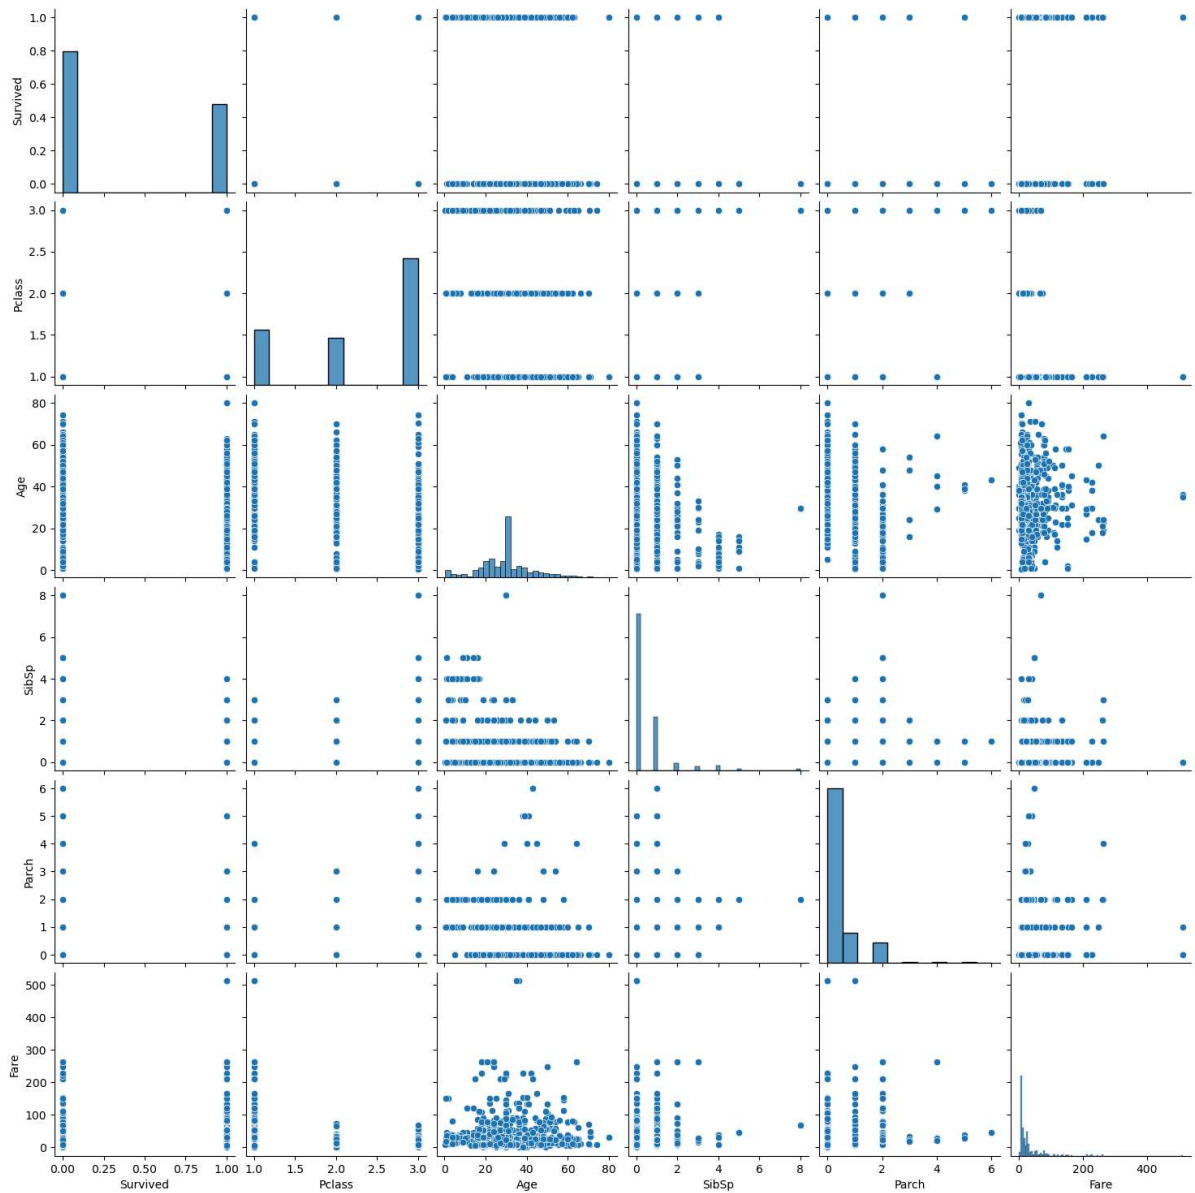
```
Out[43]:
```

	Survived	Pclass	Age	SibSp	Parch	Fare
Survived	1.000000	-0.338481	-0.069809	-0.035322	0.081629	0.257307
Pclass	-0.338481	1.000000	-0.331339	0.083081	0.018443	-0.549500
Age	-0.069809	-0.331339	1.000000	-0.232625	-0.179191	0.091566
SibSp	-0.035322	0.083081	-0.232625	1.000000	0.414838	0.159651
Parch	0.081629	0.018443	-0.179191	0.414838	1.000000	0.216225
Fare	0.257307	-0.549500	0.091566	0.159651	0.216225	1.000000

```
In [44]: # Showing Correlation Plot  
sns.heatmap(titanic.corr(),annot=True,cmap='coolwarm')  
plt.show()
```




```
In [45]: # Plotting pairplot
sns.pairplot(titanic)
plt.show()
```



```
In [46]: titanic['Survived'].value_counts()
```

```
Out[46]: 0    549
         1    342
         Name: Survived, dtype: int64
```

```
In [47]: sns.countplot(x=titanic['Survived'])  
plt.show()
```




```
In [49]: from sklearn.preprocessing import LabelEncoder
# Create an instance of LabelEncoder
le = LabelEncoder()

# Apply Label encoding to each categorical column
for column in ['Sex', 'Embarked']:
    titanic[column] = le.fit_transform(titanic[column])

titanic.head()

# Sex Column

# 0 represents female
# 1 represents Male

# Embarked Column

# 0 represents C
# 1 represents Q
# 2 represents S
```

```
Out[49]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	0	3	1	22.0	1	0	7.2500	2
1	1	1	0	38.0	1	0	71.2833	0
2	1	3	0	26.0	0	0	7.9250	2
3	1	1	0	35.0	1	0	53.1000	2
4	0	3	1	35.0	0	0	8.0500	2

```
In [50]: # importing libraries

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
```

```
In [51]: cols = ['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked']
x = titanic[cols]
y = titanic['Survived']
print(x.shape)
print(y.shape)
print(type(x)) # DataFrame
print(type(y)) # Series
```

```
(891, 7)
(891,)
<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.series.Series'>
```

In [52]: `x.head()`

Out[52]:

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	3	1	22.0	1	0	7.2500	2
1	1	0	38.0	1	0	71.2833	0
2	3	0	26.0	0	0	7.9250	2
3	1	0	35.0	1	0	53.1000	2
4	3	1	35.0	0	0	8.0500	2

In [53]: `y.head()`

Out[53]:

```
0    0
1    1
2    1
3    1
4    0
Name: Survived, dtype: int64
```

In [54]: `print(891*0.10)`

89.100000000000001

In [55]:

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.10,random_state=42)
print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(801, 7)
(90, 7)
(801,)
(90,)
```

In [56]:

```
def cls_eval(ytest,ypred):
    cm = confusion_matrix(ytest,ypred)
    print('Confusion Matrix\n',cm)
    print('Classification Report\n',classification_report(ytest,ypred))

def mscore(model):
    print('Training Score',model.score(x_train,y_train)) # Training Accuracy
    print('Testing Score',model.score(x_test,y_test))   # Testing Accuracy
```

```
In [57]: # Building the Logistic Regression Model
lr = LogisticRegression(max_iter=1000,solver='liblinear')
lr.fit(x_train,y_train)
```

Out[57]: LogisticRegression(max_iter=1000, solver='liblinear')

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [58]: # Computing Training and Testing score
mscore(lr)
```

Training Score 0.8052434456928839
Testing Score 0.7666666666666667

```
In [59]: # Generating Prediction
ypred_lr = lr.predict(x_test)
print(ypred_lr)
```

```
[1 0 1 1 1 0 0 1 0 1 0 1 0 0 1 0 0 0 0 1 0 0 1 0 1 0 1 1 0 1 1 0 0 1 0 0 0
 0 0 0 1 1 1 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 0 1 0 0 1 0 0 0 0 0
 1 0 1 0 0 1 0 0 0 0 1 0 0 0 0 1]
```

```
In [60]: # Evaluate the model - confusion matrix, classification Report, Accuracy score
cls_eval(y_test,ypred_lr)
acc_lr = accuracy_score(y_test,ypred_lr)
print('Accuracy Score',acc_lr)
```

Confusion Matrix

```
[[46  7]
```

```
[14 23]]
```

Classification Report

	precision	recall	f1-score	support
0	0.77	0.87	0.81	53
1	0.77	0.62	0.69	37
accuracy			0.77	90
macro avg	0.77	0.74	0.75	90
weighted avg	0.77	0.77	0.76	90

Accuracy Score 0.7666666666666667

```
In [61]: # Building the knnClassifier Model
knn=KNeighborsClassifier(n_neighbors=8)
knn.fit(x_train,y_train)
```

Out[61]: KNeighborsClassifier(n_neighbors=8)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [62]: # Computing Training and Testing score
mscore(knn)
```

Training Score 0.7752808988764045
Testing Score 0.6777777777777778

```
In [63]: # Generating Prediction
ypred_knn = knn.predict(x_test)
print(ypred_knn)
```

```
[1 0 0 1 1 0 0 1 1 0 0 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 1
 0 1 0 0 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 1 0 0 0 1 0
 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0]
```

```
In [64]: # Evaluate the model - confusion matrix, classification Report, Accuracy score
cls_eval(y_test,ypred_knn)
acc_knn = accuracy_score(y_test,ypred_knn)
print('Accuracy Score',acc_knn)
```

Confusion Matrix

```
[[47  6]
 [23 14]]
```

Classification Report

	precision	recall	f1-score	support
0	0.67	0.89	0.76	53
1	0.70	0.38	0.49	37
accuracy			0.68	90
macro avg	0.69	0.63	0.63	90
weighted avg	0.68	0.68	0.65	90

Accuracy Score 0.6777777777777778

```
In [65]: # Building Support Vector Classifier Model
svc = SVC(C=1.0)
svc.fit(x_train, y_train)
```

Out[65]: SVC()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [66]: # Computing Training and Testing score
mscore(svc)
```

Training Score 0.6891385767790262
Testing Score 0.6333333333333333

```
In [67]: # Generating Prediction
ypred_svc = svc.predict(x_test)
print(ypred_svc)
```

```
[0 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 1 0 0 1 0 0 0 0 0 0
 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 1 0
 0 0 0 0 0 1 1 1 0 0 0 0 0 0 0 0]
```

```
In [68]: # Evaluate the model - confusion matrix, classification Report, Accuracy score
cls_eval(y_test,ypred_svc)
acc_svc = accuracy_score(y_test,ypred_svc)
print('Accuracy Score',acc_svc)
```

Confusion Matrix

```
[[48  5]
 [28  9]]
```

Classification Report

	precision	recall	f1-score	support
0	0.63	0.91	0.74	53
1	0.64	0.24	0.35	37
accuracy			0.63	90
macro avg	0.64	0.57	0.55	90
weighted avg	0.64	0.63	0.58	90

Accuracy Score 0.6333333333333333

```
In [69]: # Building the RandomForest Classifier Model
rfc=RandomForestClassifier(n_estimators=80,criterion='entropy',min_samples_split=
rfc.fit(x_train,y_train)
```

```
Out[69]: RandomForestClassifier(criterion='entropy', max_depth=10, min_samples_split=
5,
                                n_estimators=80)
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [70]: # Computing Training and Testing score
mscore(rfc)
```

Training Score 0.9225967540574282

Testing Score 0.7666666666666667

```
In [71]: # Generating Prediction
ypred_rfc = rfc.predict(x_test)
print(ypred_rfc)
```

```
[1 0 1 1 1 0 0 1 1 1 0 0 0 0 1 1 0 0 0 1 0 0 1 0 1 0 1 0 1 1 0 1 1 0 0 0 0 0 1
 0 1 0 0 1 1 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0
 1 0 1 0 0 1 0 0 0 0 1 0 0 0 0 1]
```

```
In [72]: # Evaluate the model - confusion matrix, classification Report, Accuracy score
cls_eval(y_test,ypred_rfc)
acc_rfc = accuracy_score(y_test,ypred_rfc)
print('Accuracy Score',acc_rfc)
```

Confusion Matrix

```
[[47  6]
```

```
[15 22]]
```

Classification Report

	precision	recall	f1-score	support
0	0.76	0.89	0.82	53
1	0.79	0.59	0.68	37
accuracy			0.77	90
macro avg	0.77	0.74	0.75	90
weighted avg	0.77	0.77	0.76	90

Accuracy Score 0.7666666666666667

```
In [73]: # Building the DecisionTree Classifier Model
dt = DecisionTreeClassifier(max_depth=5,criterion='entropy',min_samples_split=1)
dt.fit(x_train, y_train)
```

Out[73]: DecisionTreeClassifier(criterion='entropy', max_depth=5, min_samples_split=1)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [74]: # Computing Training and Testing score
mscore(dt)
```

Training Score 0.8526841448189763

Testing Score 0.7777777777777778

```
In [75]: # Generating Prediction
ypred_dt = dt.predict(x_test)
print(ypred_dt)
```

```
[1 0 1 1 1 0 0 1 0 1 0 1 0 0 1 0 0 0 0 1 0 0 1 0 1 0 1 1 0 1 1 0 0 1 0 0 1
 0 0 0 1 1 1 0 1 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 1 0 0 1 0 0 0 0 0
 1 0 1 0 0 1 0 0 0 0 1 0 0 0 0 1]
```

```
In [76]: # Evaluate the model - confusion matrix, classification Report, Accuracy score
cls_eval(y_test,ypred_dt)
acc_dt = accuracy_score(y_test,ypred_dt)
print('Accuracy Score',acc_dt)
```

Confusion Matrix

```
[[46  7]
```

```
[13 24]]
```

Classification Report

	precision	recall	f1-score	support
0	0.78	0.87	0.82	53
1	0.77	0.65	0.71	37
accuracy			0.78	90
macro avg	0.78	0.76	0.76	90
weighted avg	0.78	0.78	0.77	90

Accuracy Score 0.7777777777777778

```
In [77]: # Building the Adaboost model
ada_boost = AdaBoostClassifier(n_estimators=80)
ada_boost.fit(x_train,y_train)
```

Out[77]: AdaBoostClassifier(n_estimators=80)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [78]: # Computing the Training and Testing Score
mscore(ada_boost)
```

Training Score 0.8564294631710362

Testing Score 0.7666666666666667

```
In [79]: # Generating the predictions
ypred_ada_boost = ada_boost.predict(x_test)
```

```
In [80]: # Evaluate the model - confusion matrix, classification Report, Accuracy Score
cls_eval(y_test,ypred_ada_boost)
acc_adab = accuracy_score(y_test,ypred_ada_boost)
print('Accuracy Score',acc_adab)
```

Confusion Matrix

```
[[45  8]
```

```
[13 24]]
```

Classification Report

	precision	recall	f1-score	support
0	0.78	0.85	0.81	53
1	0.75	0.65	0.70	37
accuracy			0.77	90
macro avg	0.76	0.75	0.75	90
weighted avg	0.77	0.77	0.76	90

Accuracy Score 0.7666666666666667

```
In [81]: models = pd.DataFrame({
        'Model': ['Logistic Regression','knn','SVC','Random Forest Classifier','Decision Tree Classifier'],
        'Score': [acc_lr,acc_knn,acc_svc,acc_rfc,acc_dt,acc_adab]})

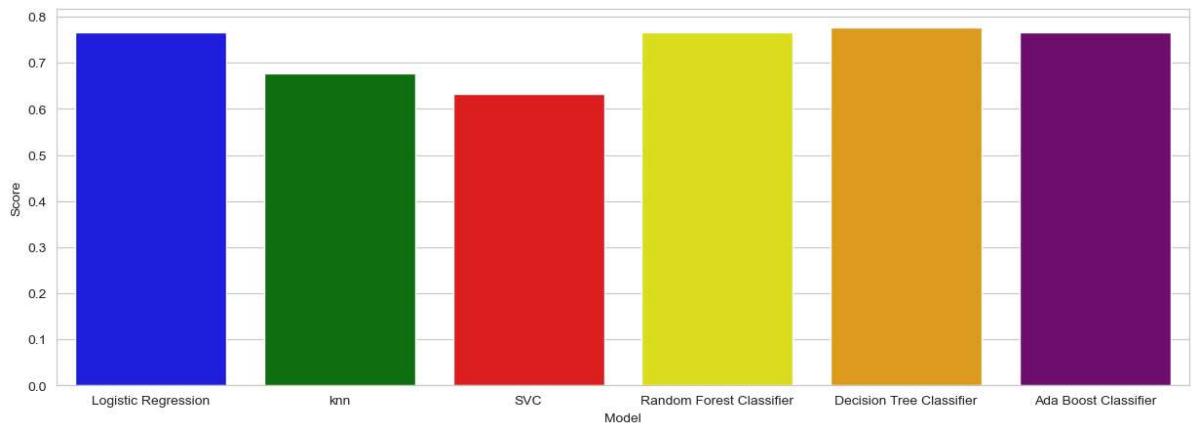
models.sort_values(by = 'Score', ascending = False)
```

Out[81]:

	Model	Score
4	Decision Tree Classifier	0.777778
0	Logistic Regression	0.766667
3	Random Forest Classifier	0.766667
5	Ada Boost Classifier	0.766667
1	knn	0.677778
2	SVC	0.633333


```
In [82]: colors = ["blue", "green", "red", "yellow", "orange", "purple"]

sns.set_style("whitegrid")
plt.figure(figsize=(15,5))
plt.ylabel("Accuracy %")
plt.xlabel("Algorithms")
sns.barplot(x=models['Model'],y=models['Score'], palette=colors )
plt.show()
```



```
In [ ]:
```